Mathematical and Architectural Feasibility of Bayesian Large Language Models for Epistemic Uncertainty Quantification

The transition of Large Language Models (LLMs) from experimental novelties to the bedrock of modern computational infrastructure has exposed a critical vulnerability: the lack of robust, principled uncertainty quantification. In high-stakes environments such as clinical decision-making, legal analysis, and autonomous system regulation, the ability of a model to express doubt is as vital as its ability to provide answers.[1] Traditional deep learning models, including standard Transformer-based LLMs, operate primarily as frequentist point-estimators. They produce a single set of weights θ that maximize a likelihood function, resulting in a predictive distribution $p(y|x, \theta)$ that often exhibits pathological overconfidence, especially when encountering out-of-distribution (OOD) inputs or ambiguous queries.[2]

Epistemic uncertainty, which represents the model's lack of knowledge due to limited data or architectural constraints, remains the most elusive component to quantify in generative systems.[1] While aleatoric uncertainty (inherent randomness in the data) can be captured through the entropy of the output logits, epistemic uncertainty requires a marginalization over the parameter space. The Bayesian framework offers a mathematically rigorous solution by treating parameters as random variables and inferring a posterior distribution $p(\theta|\mathcal{D})$. However, the sheer scale of modern LLMs—often exceeding hundreds of billions of parameters—has historically made true Bayesian inference computationally prohibitive. This report examines the mathematical and architectural feasibility of Bayesian Large Language Models (B-LLMs), analyzing the friction between autoregressive objectives and variational inference, the tractability of high-dimensional posteriors, and the state of the tooling stack in 2026.

Theoretical Foundations of Epistemic Uncertainty

The core objective of Bayesian inference in deep learning is the estimation of the posterior distribution over model weights, $p(\theta|\mathcal{D})$. This distribution encapsulates the degree of belief in different parameter configurations after observing the training data $\mathcal{D}$.[7] The predictive

distribution for a new query x^* is calculated through Bayesian Model Averaging (BMA):

$$p(y^*|x^*, \mathcal{D}) = \int p(y^*|x^*, \theta)p(\theta|\mathcal{D})d\theta$$

In practice, this integral is intractable for large-scale neural networks. Frequentist heuristics such as Monte Carlo (MC) dropout, deep ensembles, and temperature scaling have been widely adopted as "approximate Bayesian" surrogates due to their simplicity and lower computational cost.[9] However, these methods often fail to provide a true measure of epistemic uncertainty, as they do not perform principled posterior contraction in the weight space.[11]

Epistemic vs. Aleatoric Decompositions in Generative Contexts

In the context of generative LLMs, the total uncertainty of a prediction can be decomposed using information-theoretic measures. Given a query x, the total uncertainty in the predicted token y is represented by the entropy of the predictive distribution, H. This can be split as:

$$H = \mathbb{E}_{p(\theta|\mathcal{D})}[H[p(y|x, \theta)]] + I(\theta; y|x, \mathcal{D})$$

The first term on the right-hand side represents the expected aleatoric uncertainty—the average randomness in the model's predictions across the posterior. The second term, the Mutual Information (MI) between the parameters and the output, serves as a proxy for epistemic uncertainty.[4] It quantifies how much information about the parameters would be gained if the "true" next token were revealed. When epistemic uncertainty is high, different parameter samples from the posterior $p(\theta|\mathcal{D})$ will yield widely divergent predictions, signaling that the model lacks sufficient knowledge to commit to a single answer.[4]

Recent evidence suggests that LLMs exhibit a "feature gap" where the distance between the hidden representations of a standard model and an "ideal" model (one with perfect knowledge) can be used to quantify epistemic gaps.[4] This distance is hypothesized to manifest in specific semantic dimensions such as context-reliance, comprehension, and honesty.[4] For instance, a model's failure to ground an answer in the provided context (context-reliance) often points toward epistemic uncertainty rather than aleatoric noise.[4]

The Mapping Problem: Weight-Space vs. Function-Space

A central theoretical hurdle in Bayesian LLMs is the "mapping problem." This refers to the

complex, non-linear transformation that translates a distribution in the weight space Θ to a distribution in the function space $\mathcal{F}$.[13] Traditional Bayesian Neural Networks (BNNs) often employ independent Gaussian priors, $\mathcal{N}(0, \sigma^2 I)$, on individual weights. However, in the high-dimensional and highly structured architecture of a Transformer, such priors often fail to induce meaningful functional diversity.[13]

Why Weight-Space Priors Fail in Transformers

The Transformer's reliance on residual connections, layer normalization, and multiplicative attention mechanisms creates a unique geometric landscape for the posterior. In this landscape, large regions of the weight space may be functionally equivalent (symmetries), while infinitesimal movements in other directions may cause catastrophic changes in the model's output.[13]

Research indicates that standard weight-space inference, regardless of the approximation (e.g., Mean-Field Variational Inference), often struggles to surpass simple frequentist baselines in calibration.[13] The problem is attributed to the difficulty of specifying "meaningful" priors. For a Transformer to exhibit diverse functional behavior—essential for epistemic quantification—the prior must capture the intricate dependencies between the attention heads and the feed-forward blocks.[13] Failure to do so results in a posterior that collapses into a single functional mode, even if the weight-space variance remains high.[13]

To mitigate the mapping problem, researchers have proposed moving closer to the function space. Methods like Dirichlet Attention apply variational inference directly to the attention weights, treating the routing of information as a probabilistic process governed by a Dirichlet distribution.[13] This approach ensures that the uncertainty is modeled at a level that directly impacts the functional behavior of the model, bypassing some of the complexities of the full weight-space mapping.[13]

Autoregressive Objective and Variational Inference Friction

The sequential, token-by-token generation of autoregressive LLMs introduces fundamental mathematical friction with standard Variational Inference (VI). In VI, we seek to approximate the true posterior $p(\theta|\mathcal{D})$ with a simpler variational distribution $q_\phi(\theta)$ by maximizing the Evidence Lower Bound (ELBO) [8]:

$$\text{ELBO}(\phi) = \mathbb{E}{q\phi(\theta)} - D_{KL}(q_\phi(\theta) |

| p(\theta))$$

In a decoder-only LLM, the log-likelihood term, $\log p(\mathcal{D}|\theta)$, is a sum of log-conditional probabilities across thousands of tokens. The sequential nature of this objective creates a compounding effect that can interfere with the stability of the variational posterior.[16]

Variance Drift and the Curse of Recursion

A critical phenomenon in long-form Bayesian generation is "variance drift" or "variance amplification." As the model generates each new token y_t, it conditions its prediction on the previously generated tokens $y_{<t}$. If these previous tokens were sampled from a distribution with its own variance, and if the model parameters θ are also being sampled from a posterior $q_\phi(\theta)$, the uncertainty at each step compounds.[16]

Theoretical models of recursive systems suggest that without constant external grounding (new, authentic data), the model's representation of "truth" begins to drift like a random walk.[16] This is particularly problematic in "self-improving" or "self-training" loops where a model is fine-tuned on its own generated outputs. The dynamics of such systems often lead to two failure modes:

1. **Entropy Decay**: A monotonic loss of distributional diversity as the model converges on its own most likely (but potentially incorrect) outputs, a process similar to mode collapse.[16]
2. **Semantic Drift**: The model's internal beliefs move away from the true distribution, bounded only by the support diameter of the latent space.[16]

In long-form Bayesian generation, this drift can make the output increasingly incoherent. As the sequence length increases, the cumulative variance in the parameter space may lead to tokens that are individually plausible but collectively nonsensical, as the "model variant" represented by a single sample of θ may not be stable across a 1,000-token generation horizon.[16]

The "Equivalent Heuristic" Hypothesis: MC-Dropout vs. Deep Ensembles

A long-standing debate in the field is whether advanced Bayesian methods provide a non-trivial improvement over frequentist heuristics. The "Equivalent Heuristic" hypothesis suggests that methods like Monte Carlo (MC) Dropout or Deep Ensembles might be mathematically equivalent to specific classes of Bayesian models in the limit of infinite

samples or specific architecture widths.[1]

Posterior Contraction and the Bayesian Advantage

Mathematical analysis reveals that while heuristics may mimic some aspects of uncertainty, they lack the "posterior contraction" property inherent in true Bayesian inference.[1] In a Bayesian setting, as the volume of training data $\mathcal{D}$ increases, the posterior $p(\theta|\mathcal{D})$ should contract toward the "true" parameter configuration.

- **MC-Dropout** is essentially an ensemble of shared-weight models with a fixed noise mask. It does not possess a mechanism to update its "prior" belief based on data; the dropout rate is typically a fixed hyperparameter rather than a learned parameter of a posterior.[1]
- **Deep Ensembles** have been shown to be highly effective at capturing different modes of the posterior, often outperforming single-mode Bayesian approximations like Mean-Field VI.[12] However, they are not a true distribution; they are a set of point estimates. They lack the ability to perform Bayesian Model Averaging (BMA) over the continuous density of the parameter space.[1]

Empirical benchmarks on LLMs indicate that "straightforward combinations" of parameter-efficient fine-tuning (PEFT) and basic Bayesian techniques like MC-Dropout or Ensembles yield only marginal improvements in generalization compared to structured methods like BLoB (Bayesian Low-Rank Adaptation by Backpropagation).[11] BLoB and its successors provide a unique "second-order" uncertainty that reflects the curvature of the loss landscape, a feature that frequentist heuristics only capture indirectly.[12]

Architectural Specificity: Bayesian Treatment of Transformer Components

The feasibility of B-LLMs depends heavily on which parts of the Transformer block are chosen for Bayesian inference. A typical Transformer layer consists of the Multi-Head Attention (MHA) block and the Feed-Forward Network (FFN).[13]

FFN as the Knowledge Store: The Primary Bayesian Candidate

Research using "Bayesian Wind Tunnels"—synthetic environments where the analytic posterior is known—has provided mechanistic insights into how Transformers handle uncertainty.[15] These studies have shown that the FFNs are responsible for the numerical belief updates.[15]

- **Residual Stream as Belief Substrate**: The residual stream acts as a coordinate system where hypotheses and their probabilities are tracked layer-by-layer.[15]
- **FFNs as Update Engines**: FFNs perform the heavy lifting of numeric belief refinement.

Probes of the delta-mean squared error show that FFNs steadily improve the model's belief readout across the depth of the network.[24]

- **Attention as Routing**: The Multi-Head Attention mechanism serves primarily to route information. It determines *which* components of the belief substrate should be updated by the FFNs but does not independently drive the precision of the posterior belief.[15]

Given this division of labor, the FFN blocks (specifically the second linear layer of each block) are the most effective candidates for Bayesian treatment. This is particularly relevant in Mixture-of-Experts (MoE) architectures, where applying structured Laplace approximations to the expert pathways allows for calibrated uncertainty estimation without the need to "Bayesianize" the entire 100B+ parameter backbone.[25]

Non-Transformer Alternatives: State Space Models (SSMs)

The quadratic complexity of the Attention mechanism has led to the rise of alternative sequence models like Mamba and Hyena. These models are naturally "Bayesian-friendly" due to their recursive, linear structure.[15]

- **Recursive State Tracking**: Mamba's internal state acts similarly to a hidden state in a Kalman filter or a Hidden Markov Model (HMM).[15] This makes the mathematical integration of Bayesian filtering (e.g., Extended Kalman Filters) far more direct than in the global, non-recursive attention of Transformers.[26]
- **Infinite-Dimensional Mappings**: Models like DeepONet and Neural Operators are designed to learn mappings between function spaces rather than just point-to-point discrete values.[27] These architectures may ultimately prove superior for epistemic quantification as they treat the model's output as a continuous field of probabilities rather than a discrete sequence of tokens.[27]

Tractability of p(theta | D) for N > 1B Parameters

A common criticism of Bayesian methods is the lack of tractability for models with billions of parameters. However, recent advances have provided a mathematical path forward, primarily through subspace inference and marginalization.[18]

The Subspace Proof of Tractability

For a model with N parameters, the full posterior $p(\theta|\mathcal{D})$ is indeed intractable due to the $O(N^2)$ size of the covariance matrix. However, the tractability can be proven for a restricted subspace $\Theta_{sub} \subset \Theta$.

1. **Decomposition**: Using Low-Rank Adaptation (LoRA), we define the parameter update as $\Delta W = BA$, where $B \in \mathbb{R}^{m \times r}$ and $A \in \mathbb{R}^{r \times n}$.[12]

2. **Marginalization**: The full posterior is approximated by a marginal distribution over the low-rank matrices A and B, while keeping the pre-trained weights W_0 fixed.[18]
3. **Complexity Reduction**: For a rank $r \ll \min(m, n)$, the number of trainable parameters k is $r(m+n)$, which is typically less than 1% of the total parameter count.[18]
4. **Tractable Inference**: The posterior $p(A, B|\mathcal{D})$ can be estimated using structured Laplace approximations or Variational Inference (e.g., BLoB). For $k \approx 10^6$ parameters, the $O(k^2)$ Hessian or $O(k)$ diagonal approximation is computationally feasible on 2026 hardware accelerators.[9]

This proof shows that B-LLMs do not require the Bayesianization of the entire weight space to provide meaningful epistemic uncertainty; capturing the uncertainty in the "adaptation space" is sufficient for significant calibration gains.[9]

Tooling and Library State-of-the-Art in 2026

The implementation of B-LLMs relies on a specialized "Bayesian Stack." A comparative analysis of current research and industry tools reveals a clear divergence based on performance and ease of use.[23]

The JAX vs. PyTorch Ecosystem

Framework	Primary Bayesian Libraries	Strengths	Best For
JAX	BlackJAX, Distrax	High-performance XLA loops, auto-vectorized sampling, functional purity.[23]	High-speed MCMC research, complex posterior samplers.[23]
PyTorch	Laplace, Pyro, ASDL	Industry standard for LLMs, native integration with Hugging Face/PEFT.[9]	Post-hoc Laplace approximations, Variational LoRA, production deployment.[9]

TensorFlow	TF Probability (TFP)	Legacy viability, robust distribution library.[23]	Continued support for existing 2020-era medical AI pipelines.[23]

PyTorch + Laplace/ASDL

The Laplace library has become the industry standard for post-hoc Bayesianization. It allows researchers to take a standard, fine-tuned LoRA model and approximate its posterior by calculating the Hessian of the loss at the MAP (Maximum A Posteriori) point.[9] When combined with the ASDL (Automatic Second-Order Differentiation Library), it enables Kronecker-factorized approximations that balance accuracy and speed.[23]

JAX + BlackJAX/Distrax

For researchers pushing the boundaries of Stochastic Gradient MCMC (SG-MCMC), JAX is the preferred environment. BlackJAX provides a library of advanced samplers (HMC, NUTS) that leverage JAX's just-in-time compilation to run thousands of parallel chains, which is essential for verifying the convergence of the posterior in deep models.[23]

The Status of TensorFlow Probability (TFP)

While TensorFlow usage has shifted in the LLM era, TFP remains viable for its unparalleled library of probability distributions and bijectors. In 2026 hardware contexts, TFP's integration with XLA allows it to remain competitive for smaller-scale Bayesian modeling and ensemble management, though it is rarely the first choice for billion-parameter Transformers.[23]

Efficiency Techniques: Laplace LoRA vs. SG-MCMC vs. VI

Quantifying the trade-offs between different Bayesian approximations is essential for selecting the right "stack" for a specific project.

Laplace Approximation on LoRA (Post-hoc)

Laplace-LoRA is characterized by its post-hoc nature; the model is fine-tuned normally, and the Bayesian treatment is applied as a final step.[5]

- **Pros**: No changes to the training pipeline; minimal compute overhead (+1-10% latency); extremely memory efficient.[9]
- **Cons**: Relies on a Gaussian approximation around a single mode; can be inaccurate if the loss landscape is highly non-Gaussian.[12]

Variational Inference (IVON / BLoB)

Variational methods like IVON and BLoB learn the mean and covariance simultaneously during training.[5]

- **Pros**: Superior calibration; can capture "non-local" uncertainty better than Laplace; IVON acts as a drop-in replacement for AdamW.[29]
- **Cons**: Increases training parameters (often by 2x for the diagonal variance); requires more tuning of the prior precision and KL-scaling factor.[18]

Stochastic Gradient MCMC (SG-MCMC)

SG-MCMC provides a collection of samples that represent the true posterior without Gaussian assumptions.

- **Pros**: Theoretically the most accurate representation of the posterior; handles multi-modality naturally.[28]
- **Cons**: Extremely slow; requires storing multiple full-rank parameter sets; typically only tractable for the smallest LLM adapters.[12]

Survey Matrix: Comparison of Bayesian Techniques for LLMs

The following matrix provides a quantitative and qualitative comparison of the primary methods for uncertainty quantification in the current landscape.[5]

Method	Latency (Inference)	Memory Overhead	Calibration Error (ECE)	Robustness (OOD)
MAP LoRA	Baseline (1x)	0%	High (Poor)	Low
MC-Dropout	N Forward Passes	Modest	Medium	Medium
Deep Ensembles	N Forward Passes	Very High	High	High
Laplace-LoRA	1.1x	+1-5%	Low	High
BLoB (Variational)	N Forward Passes	+33% to +50%	Very Low	Very High
TFB	1.2x	+5%	Low	High

(Training-Free)				
IVON-LoRA	1x (Mean Pred)	Minimal	Low	Medium
Me-LoRA	N Forward Passes	+10%	Very Low	Very High

Note: ECE (Expected Calibration Error) and robustness scores are relative to a standard fine-tuned Llama-2 or Llama-3 baseline.[2]

Task-Specific Utility: Hallucinations and OOD Detection

The primary motivation for implementing B-LLMs is the resolution of task-specific failures that plague deterministic models.

Hallucination Attribution

Bayesian posteriors provide a principled way to distinguish between "I don't know because the data is noisy" (aleatoric) and "I don't know because I haven't seen this" (epistemic).[2]

- **Epistemic Attribution**: By measuring the variance of the predictions across posterior samples, a system can flag a "knowledge gap." If different weight samples produce radically different "facts," the model is likely hallucinating due to a lack of training data.[2]
- **Aleatoric Attribution**: If the model is certain in its *uncertainty* (i.e., the parameter samples agree that the output distribution should be uniform), the system has identified a contradictory or ambiguous prompt.[2]

Recent ABC (Approximate Bayesian Computation) approaches treat the LLM as a "stochastic simulator" to infer posterior distributions over predictive probabilities, reducing Brier scores by up to 74% and significantly enhancing calibration in clinical diagnostics.[1]

Out-of-Distribution (OOD) Detection

Deterministic LLMs are notorious for "confidently making things up" when presented with inputs outside their training domain.[2] B-LLMs show significantly better OOD performance.

- **Benchmark Results**: On tasks involving code in non-existent languages or rare biomedical interactions, B-LLMs like BLoB and TFB maintain low confidence, whereas standard LoRA models often produce high-confidence, syntactically correct but semantically void output.[5]

- **Mechanism**: The posterior $p(\theta|\mathcal{D})$ represents the "volume" of parameters consistent with the training data. When an input x^* is OOD, it typically activates features that were not constrained during training, leading to a massive divergence in the predictions of different posterior samples.[7]

Framework Selection: A Decision Tree for 2026 Implementation

Choosing the right framework and approximation method depends on the project's specific constraints regarding compute, memory, and required uncertainty precision.

1. **Is the model already fine-tuned?**
   - **Yes**: Use **Training-Free Bayesianization (TFB)** if an "anchor dataset" (subset of training data) is available to search for variance scales.[5] If not, use **Laplace-LoRA** for a post-hoc Gaussian approximation.[9]
   - **No**: Proceed to the training phase.
2. **Is training compute/memory severely limited?**
   - **Yes**: Use the **IVON Optimizer**. It replaces AdamW and learns a diagonal Gaussian posterior at essentially the same cost, providing uncertainty "for free".[29]
   - **No**: Proceed to high-precision options.
3. **Is the highest possible calibration required for high-stakes decision-making?**
   - **Yes**: Use **BLoB (Bayesian LoRA by Backpropagation)** or **Me-LoRA** to jointly estimate mean and covariance during training. This provides the most robust epistemic quantification.[5]
   - **No**: Use **Deep Ensembles of LoRA adapters**, which are easier to implement if you have the storage for multiple small matrix sets.[9]
4. **Are you developing new Bayesian algorithms?**
   - **Yes**: Use the **JAX Stack (BlackJAX/Distrax)** for its flexibility in sampling and auto-vectorization.[23]
   - **No**: Use the **PyTorch Stack (Laplace/ASDL)** for its integration with the wider LLM ecosystem (Hugging Face, PEFT, DeepSpeed).[9]

Conclusions and Practical Implications

The mathematical and architectural feasibility of Bayesian Large Language Models is no longer in question; rather, the focus has shifted to choosing the right approximation for the right scale. Weight-space Bayesian inference in the low-rank subspace (LoRA) provides a non-trivial and measurable improvement over frequentist heuristics like MC-dropout, particularly in its ability to contract the posterior around the data.[9]

The "mapping problem" remains a significant challenge, but focusing Bayesian updates on the

Feed-Forward Network (FFN) blocks—the Transformer's primary numeric belief engines—offers a clear path to high-quality epistemic quantification.[13] Furthermore, the introduction of "Uncertainty Distillation" allows for the deployment of these capabilities in real-time, low-latency environments by distilling the confidence of a Bayesian "Teacher" into a non-Bayesian "Student" model.[12]

As we move into 2026, the integration of Bayesian principles into the Transformer architecture—whether through post-hoc Laplace approximations, variational optimizers like IVON, or recursive state-tracking in SSMs—will be the defining factor in creating AI systems that are not just fluent, but reliably honest about the limits of their own knowledge.[2]

Works cited

1. Uncertainty Quantification of Large Language Models using ... - arXiv, accessed January 27, 2026, https://arxiv.org/abs/2509.19375
2. Uncertainty Quantification of Large Language Models using Approximate Bayesian Computation - arXiv, accessed January 27, 2026, https://arxiv.org/html/2509.19375v1
3. Entropy-Constrained Embeddings for Safe and Reliable Enterprise Large Language Models, accessed January 27, 2026, https://www.researchgate.net/publication/396335484_Entropy-Constrained_Embeddings_for_Safe_and_Reliable_Enterprise_Large_Language_Models
4. Epistemic Uncertainty Quantification of LLMs in Contextual Question-Answering - arXiv, accessed January 27, 2026, https://www.arxiv.org/pdf/2510.02671
5. Training-Free Bayesianization for Low-Rank Adapters ... - Hao Wang, accessed January 27, 2026, http://www.wanghao.in/paper/NIPS25_TFB.pdf
6. [2503.15850] Uncertainty Quantification and Confidence Calibration in Large Language Models: A Survey - arXiv, accessed January 27, 2026, https://arxiv.org/abs/2503.15850
7. Posterior probability - Wikipedia, accessed January 27, 2026, https://en.wikipedia.org/wiki/Posterior_probability
8. Bayesian Posterior Sampling (1) Introduction - Yuanzhi Zhu, accessed January 27, 2026, https://yuanzhi-zhu.github.io/2023/05/25/Bayesian-Posterior-Sampling/
9. Bayesian-LoRA: Uncertainty in LLM Adaptation - Emergent Mind, accessed January 27, 2026, https://www.emergentmind.com/topics/bayesian-lora
10. Uncertainty Quantification with Pre-trained Language Models: A Large-Scale Empirical Analysis - ResearchGate, accessed January 27, 2026, https://www.researchgate.net/publication/372919451_Uncertainty_Quantification_with_Pre-trained_Language_Models_A_Large-Scale_Empirical_Analysis
11. BLoB: Bayesian Low-Rank Adaptation by Backpropagation for ..., accessed January 27, 2026, https://arxiv.org/abs/2406.11675
12. BLoB: Bayesian Low-Rank Adaptation by Backpropagation for Large Language Models | Request PDF - ResearchGate, accessed January 27, 2026, https://www.researchgate.net/publication/397196367_BLoB_Bayesian_Low-Rank_Adaptation_by_Backpropagation_for_Large_Language_Models

13. The Bayesian Transformer - Research Collection, accessed January 27, 2026, https://www.research-collection.ethz.ch/bitstreams/78589e06-71ce-4f6a-9fae-3c9e5b99184a/download
14. Bayesian Enhancement Models for One-to-Many Mapping in Image... - OpenReview, accessed January 27, 2026, https://openreview.net/forum?id=jJJOoLVAEm
15. The Bayesian Geometry of Transformer Attention - arXiv, accessed January 27, 2026, https://arxiv.org/html/2512.22471v2
16. (PDF) On the Limits of Self-Improving in LLMs and Why AGI, ASI and ..., accessed January 27, 2026, https://www.researchgate.net/publication/399666570_On_the_Limits_of_Self-Improving_in_LLMs_and_Why_AGI_ASI_and_the_Singularity_Are_Not_Near_Without_Symbolic_Model_Synthesis
17. Uncertainty-Aware Self-Attention Model for Time Series Prediction with Missing Values, accessed January 27, 2026, https://www.mdpi.com/2504-3110/9/3/181
18. ME-LORA: MEMORY-EFFICIENT BAYESIAN LOW- RANK ADAPTATION FOR LARGE LANGUAGE MODELS - OpenReview, accessed January 27, 2026, https://openreview.net/pdf/b2f149ec62a68d354760ca1fd47732a819ac8036.pdf
19. MMPD: Diverse Time Series Forecasting via Multi-Mode Patch Diffusion Loss - OpenReview, accessed January 27, 2026, https://openreview.net/forum?id=NEUgHT8dvH
20. AMMcs2013 - Wilfrid Laurier University, accessed January 27, 2026, https://ammcs.wlu.ca/2013/AMMCS-2013-Book-of-Abstracts.pdf
21. On the Limits of Self-Improving in LLMs and Why AGI, ASI and the Singularity Are Not Near Without Symbolic Model Synthesis - ResearchGate, accessed January 27, 2026, https://www.researchgate.net/publication/399928063_On_the_Limits_of_Self-Improving_in_LLMs_and_Why_AGI_ASI_and_the_Singularity_Are_Not_Near_Without_Symbolic_Model_Synthesis
22. Ensembling Pruned Attention Heads For Uncertainty-Aware Efficient Transformers - arXiv, accessed January 27, 2026, https://arxiv.org/html/2510.18358v1
23. Wang-ML-Lab/bayesian-peft: Bayesian Low-Rank ... - GitHub, accessed January 27, 2026, https://github.com/Wang-ML-Lab/bayesian-peft
24. How Transformers Do Bayesian Inference - YouTube, accessed January 27, 2026, https://www.youtube.com/watch?v=8MvVkbI0jtE
25. Bayesian Mixture of Experts For Large Language Models - arXiv, accessed January 27, 2026, https://arxiv.org/html/2511.08968v1
26. FastSLAM: An Efficient Solution to the Simultaneous Localization And Mapping Problem with Unknown Data Association - Sebastian Thrun, accessed January 27, 2026, http://robots.stanford.edu/papers/Thrun03g.pdf
27. Physics-informed deep operator network for traffic state estimation - arXiv, accessed January 27, 2026, https://arxiv.org/html/2508.12593v1
28. Posterior marginalization accelerates Bayesian inference for dynamical models of biological processes - PMC, accessed January 27, 2026,

https://pmc.ncbi.nlm.nih.gov/articles/PMC10589897/
29. Improving LoRA with Variational Learning | OpenReview, accessed January 27, 2026, https://openreview.net/forum?id=oY4CQ2GJkC
30. Uncertainty-Aware Adaptation of Large Language Models for Protein-Protein Interaction Analysis - ResearchGate, accessed January 27, 2026, https://www.researchgate.net/publication/388883734_Uncertainty-Aware_Adaptation_of_Large_Language_Models_for_Protein-Protein_Interaction_Analysis
31. Scalable Bayesian neural networks - Trung Trinh, accessed January 27, 2026, https://trungtr.com/assets/pdf/master_thesis.pdf
32. Efficient Uncertainty Estimation via Distillation of Bayesian Large Language Models - arXiv, accessed January 27, 2026, https://arxiv.org/html/2505.11731v1